



MAX PLANCK INSTITUTE FOR **BIOLOGY OF AGEING**



Help yourself on a remote server

bioinformatics@age.mpg.de

<https://mpg-age-bioinformatics.github.io>

Outline

- Connecting to a remote server
- Copying files to/from a remote server
- The modules system
- Shifter/Docker for HPC
- Installing software without `su` access
- Installing R packages
- Installing python packages
- Installing perl packages
- SLURM

Connecting to a remote server

Connecting to a remote server over a Secure Shell (ssh) – ssh username@remote.adress :

```
ssh JBoucas@amalia.age.mpg.de
```

With X forwarding:

```
ssh -X JBoucas@amalia.age.mpg.de
```

Download rc files for your user into your home folder in amalia

```
cd

wget https://raw.githubusercontent.com/mpg-age-
bioinformatics/cluster_first_steps/master/.bashrc

wget https://raw.githubusercontent.com/mpg-age-
bioinformatics/cluster_first_steps/master/.bash_profile

source ~/.bash_profile
source ~/.bashrc

# remember that next time you login this files will be automatically
sourced
```

Copying files to/from a remote server

Copying files over ssh to your home folder on a remote server:

```
scp file.txt UName@ServerAddress:~/
```

Copying files over ssh from your home folder on a remote server:

```
scp UName@ServerAddress:~/file.txt .
```

Both `scp` will only allow you to copy files (not directories) unless you use the `'-r` argument for `recursively`. For speed use `'-o Cipher=arcfour`.

rsync to a remote server

```
rsync -rtvh -e "ssh -c arcfour" source_folder \
UName@ServerAddress:destination
```

Not all servers will have the less costly cipher arcfour encryption algorithm and therefore you can remove the whole `'-e "ssh -c arcfour"` block.

The modules system

A centralized software system.

The modules system **loads software** (version of choice) and changes **environment variables** (eg. LD_LIBRARY_PATH)

```
module avail          # shows available modules
module whatis SAMtools # shows a description of the SAMtools module
module show SAMtools   # shows environment changes for SAMtools
module load SAMtools    # loads SAMtools
module list           # lists all loaded modules
module unload SAMtools # unloads the SAMtools module
module purge          # unloads all loaded modules
```

more on <http://modules.sourceforge.net>

Shifter/Docker for HPC

Mobile, ready to go, reproducible science!

Computers are not all the same



Created by Mike O'Brien
from Noun Project

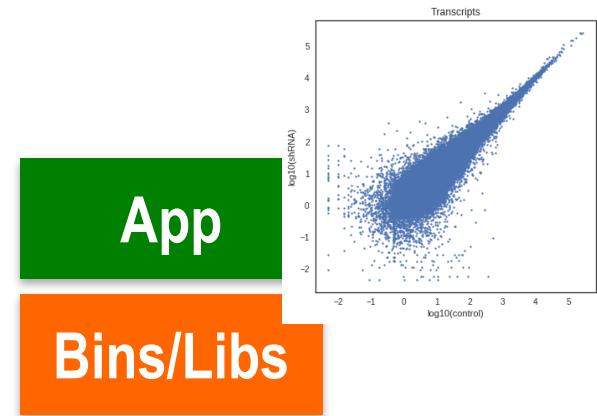
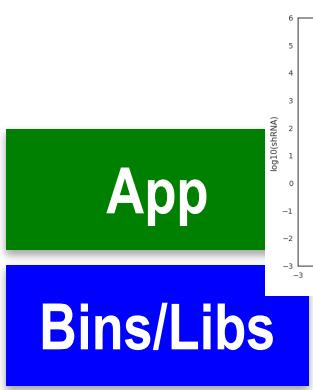


Created by Mister Pixel
from Noun Project



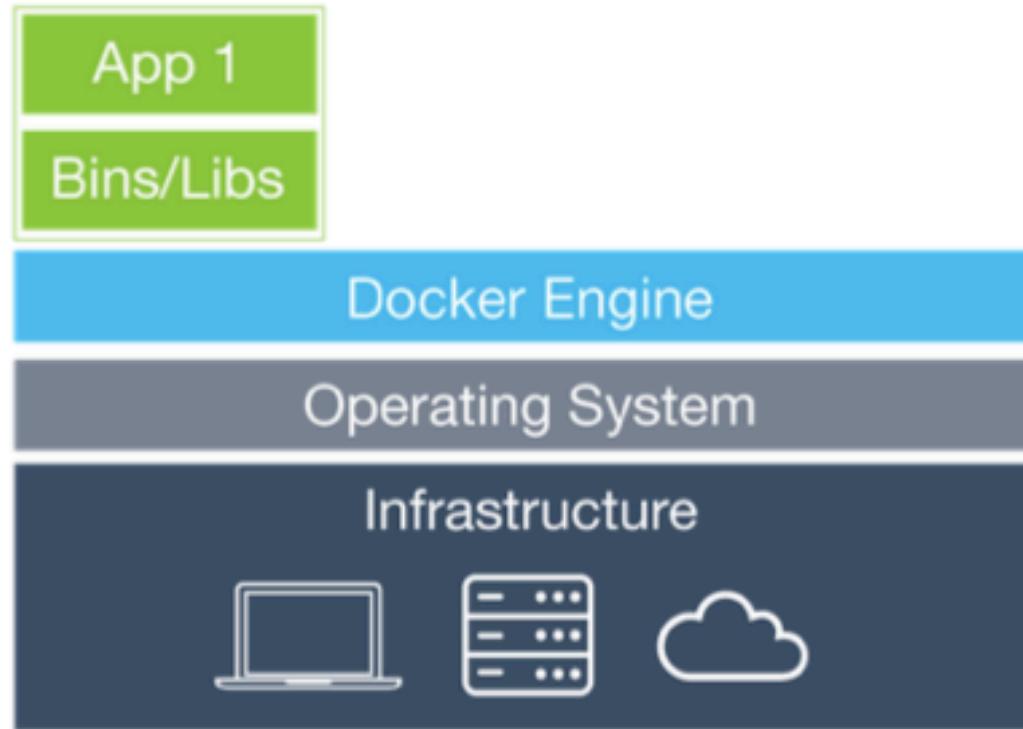
Created by factor[e] design initiative
from Noun Project

The same apparent code can give different results



Different version of dependency in bin/
== different input in one step
== code breaks

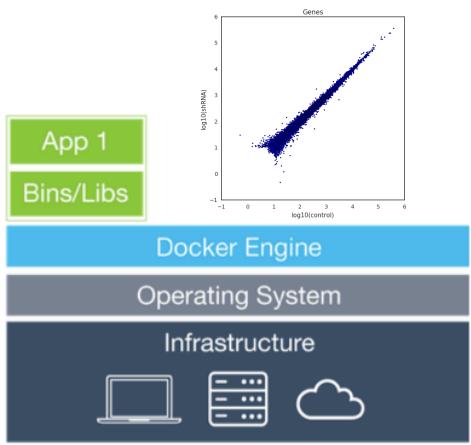
Containers



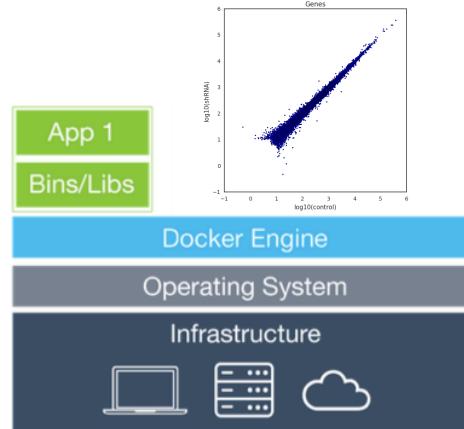
Containers

<https://blog.logentries.com/2015/07/an-all-inclusive-log-monitoring-container-for-docker/>

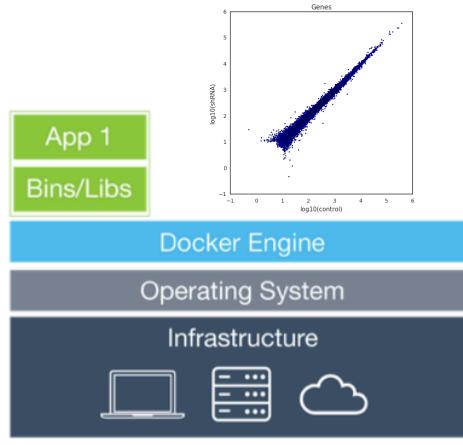
Containers are independent of their host



Containers

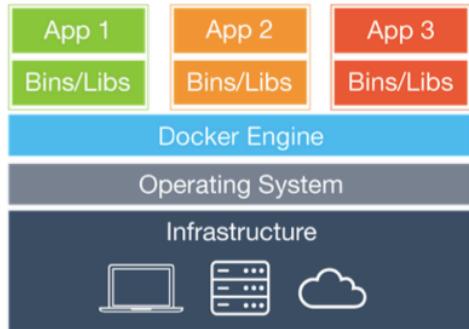


Containers

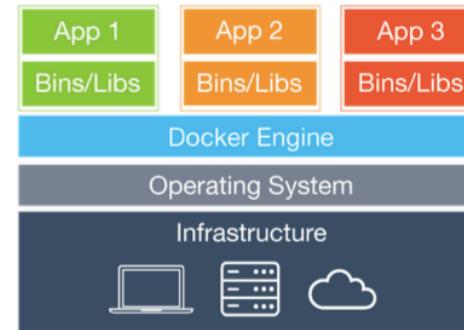


Containers

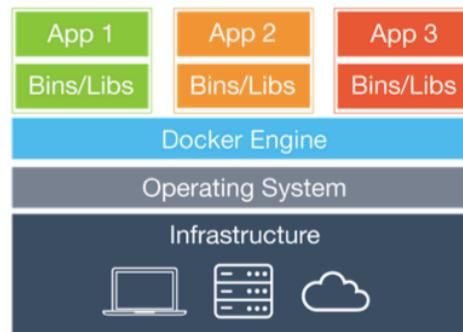
Multiple environments per machine



Containers



Containers



Containers

Shifter/Docker for HPC

```
# load respective module
module load shifter

# list available images
shifterls

# get an interactive shell on ubuntu and check you
# are in the intended image
shifter --image=ubuntu:15.10 bash -login
cat /etc/lsb-release

# use the Bioinformatics image
shifter --image=mpgabioinformatics/bioinformatics_software:v1.0.0

# run a script inside an image
shifter --image=<provider>/<image_name>:<tag> /path/to/script
```

Working interactively on the mpi age bioinformatics software image

```
module load shifter # load respective module

shifterls # list available images

which R # check if you have R

# login to our software image
shifter --image=mpgabioinformatics/bioinformatics_software:v1.0.1

pwd

module avail

module load rlang

which R
```

Running a script inside the mpi age bioinformatics software image

Example `~/test.shifter.sh`:

```
#!/bin/bash
source ~/.bashrc

module load rlang python

which python

python << EOF
print "This is python"
EOF

which R

Rscript -e "print('This is R')"
```

Running the script:

```
chmod +x test.shifter.sh
shifter \
    --image=mpgabioinformatics/bioinformatics_software:v1.0.1 \
    ./test.sh
```

Shifter/Docker for HPC

More infos on

shifter,

docker (eg. for your own laptop),

singularity (running @ draco.mpcdf.mpg.de),

and creation of images on

<http://bioinformatics.age.mpg.de/#callout> > Jupyter on Docker

https://github.com/mpg-age-bioinformatics/cluster_first_steps#shifter

More infos on our software image and how to run it on your laptop:

https://github.com/mpg-age-bioinformatics/software_docker#software-container

module avail

Example `module avail` output.

```
----- /beegfs/common/software/2017/modules/modulefiles/general -----
gcc/4.9.2          jupyterhub/0.7.2(default)    rlang/3.3.2(default)
gcc/6.3.0(default) perl/5.24.1(default)      ruby/2.4.0(default)
java/8.0.111(default) pigz/2.3.4(default)    ruby-install/0.6.1(default)
jdk/8u121(default) python/2.7.12(default)   tmux/2.3(default)
jup/0.1(default)   python/3.6.0           tools/0.1(default)

----- /beegfs/common/software/2017/modules/modulefiles/bioinformatics -----
allpathsdlg/52488(default)  gatk/3.4.46(default)    seqemehl/0.2.0(default)
bamutil/1.0.13(default)    graphviz/2.40.1(default) seqtk/1.2.r94(default)
bcl2fastq/2.17.1.14(default) hisat/2.0.4(default) skewer/0.2.2(default)
bedtools/2.24.0           igvtools/2.3.89(default) snpeff/4.3.i(default)
bedtools/2.26.0(default)   imtornado/2.0.3.3(default) spades/3.10.0(default)
bowtie/1.2.0              lofreq/2.1.2(default)    sratoolkit/2.8.1
bowtie/2.2.9(default)     meme/4.11.3           sratoolkit/2.8.1-3(default)
bwa/0.7.15(default)       meme/4.12.0(default)    star/2.5.2b(default)
bwtool/git170623(default) methpipe/3.4.2(default) stringtie/1.3.0(default)
cufflinks/2.2.1(default)  ngsutils/0.5.9(default) tophat/2.1.1(default)
cutadapt/1.13.0(default)  picard/2.8.1(default)   vcftools/0.1.14(default)
epiteome/1.0.0(default)   qiime/1.9.1(default)    walt/1.0.0(default)
expat/2.2.0(default)     rsem/1.3.0(default)
fastqc/0.11.5(default)   samtools/1.3.1(default)
```

Notice the 2 arrows and the two distinct blocks of software being shown.

The distinct blocks are generated by the different paths kept on the `MODULEPATH` variable set like this:

```
export MODULEPATH=/beegfs/common/software/2017/modules/modulefiles\
/general:/beegfs/common/software/2017/modules/modulefiles/bioinformatics
```

The modules system

`module avail` simply lists the contents of the `MODULEPATH` variable.

```
amaliax:~$ tree -L 2 /beegfs/common/software/2017/modules/modulefiles/general  
/beegfs/common/software/2017/modules/modulefiles/general  
├── gcc  
│   ├── 4.9.2  
│   └── 6.3.0  
├── java  
│   └── 8.0.111  
├── jdk  
│   └── 8u121  
├── jup  
│   └── 0.1  
├── jupyterhub  
│   └── 0.7.2  
├── perl  
│   └── 5.24.1  
├── pigz  
│   └── 2.3.4  
├── python  
│   ├── 2.7.12  
│   └── 3.6.0  
├── rlang  
│   └── 3.3.2  
├── ruby  
│   └── 2.4.0  
├── ruby-install  
│   └── 0.6.1  
├── tmux  
│   └── 2.3  
└── tools  
    └── 0.1  
  
13 directories, 15 files  
amaliax:~$ ls -la /beegfs/common/software/2017/modules/modulefiles/general/python  
total 8  
drwxrwsr-x  2 DRosskopp group_beewsw  3 Jun 24 05:15 .  
drwxrwsr-x 15 DRosskopp group_beewsw 13 Jun 24 05:15 ..  
-rw-rw-r--  1 DRosskopp group_beewsw 3078 Jun 24 05:15 2.7.12  
-rw-rw-r--  1 DRosskopp group_beewsw 3058 Jun 24 05:15 3.6.0  
-rw-rw-r--  1 DRosskopp group_beewsw 137 Jun 24 05:15 .version
```

Example content of a folder in the `MODULEPATH`.

Example content for Python.

The modules system

```
amalias:~$ ls -la /beegfs/common/software/2017/modules/modulefiles/general/python
total 8
drwxrwsr-x  2 DRoskopp group_beessw  3 Jun 24 05:15 .
drwxrwsr-x 15 DRoskopp group_beessw 13 Jun 24 05:15 ..
-rw-rw-r--  1 DRoskopp group_beessw 3078 Jun 24 05:15 2.7.12
-rw-rw-r--  1 DRoskopp group_beessw 3058 Jun 24 05:15 3.6.0
-rw-rw-r--  1 DRoskopp group_beessw 137 Jun 24 05:15 .version
```

`2.7.12` module file for python 2.7.12

`3.6.0` module file for python 3.6.0

`.version` file defining the default version:

```
##%Module1.0#####
##
## version file for python
##
set ModulesVersion    "2.7.12"
```

module file

```
##&Module3.2.10#####
proc ModulesHelp { } {
    global version
    puts stderr "\n\tVersion 2.7.12 of python\n"
}

module-whatis "Version 2.7.12 of python"

# for Tcl script use only
set      version      "3.2.10"

conflict python
prepend-path PATH /beegfs/common/software/2017/modules/software/python/2.7.12/bin
prepend-path LD_LIBRARY_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/lib
prepend-path CPATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path C_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path CPLUS_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path OJJC_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path MANPATH /beegfs/common/software/2017/modules/software/python/2.7.12/share/man
prepend-path INFODIR /beegfs/common/software/2017/modules/software/python/2.7.12/share/man
set home $::env(HOME)
set pythonuser $home/.python/2.7.12/bin
exec /bin/mkdir -p $pythonuser
prepend-path PATH $home/.python/2.7.12/bin
set jupyter_runtime_dir $home/.python/2.7.12/jupyter/run
exec /bin/mkdir -p $jupyter_runtime_dir
setenv JUPYTER_RUNTIME_DIR $home/.python/2.7.12/jupyter/run
set jupyter_data_dir $home/.python/2.7.12/jupyter/data
exec /bin/mkdir -p $jupyter_data_dir
setenv JUPYTER_DATA_DIR $home/.python/2.7.12/jupyter/data
setenv PYTHONHOME /beegfs/common/software/2017/modules/software/python/2.7.12/
setenv PYTHONPATH /beegfs/common/software/2017/modules/software/python/2.7.12/lib/python2.7
setenv PYTHONUSERBASE $home/.python/2.7.12/
exec /bin/mkdir -p $home/.python/2.7.12/pythonpath/site-packages
module load gcc/6.3.0 bzip2/1.0.6 xz/5.2.2 ncurses/6.0 libevent/2.0.22 pcre/8.39 curl/7.51.0 freetype/2.7 openblas/0.2.19
setenv CFLAGS "-I/beegfs/common/software/2017/modules/software/openblas/0.2.19/include -I/beegfs/common/software/2017/modules/software/ncurses/6.0/include/ncurses -I/beegfs/common/software/2017/modules/software/libevent/2.0.22/include -I/beegfs/common/software/2017/modules/software/bzip2/1.0.6/include -I/beegfs/common/software/2017/modules/software/xz/5.2.2/include -I/beegfs/common/software/2017/modules/software/pcre/8.39/include -I/beegfs/common/software/2017/modules/software/curl/7.51.0/include -I/beegfs/common/software/2017/modules/software/openblas/0.2.19/include -I/beegfs/common/software/2017/modules/software/rlang/3.3.2/lib64/R/include"
setenv LDFLAGS "-L/beegfs/common/software/2017/modules/software/openblas/0.2.19/lib -L/beegfs/common/software/2017/modules/software/ncurses/6.0/lib -L/beegfs/common/software/2017/modules/software/libevent/2.0.22/lib -L/beegfs/common/software/2017/modules/software/bzip2/1.0.6/lib -L/beegfs/common/software/2017/modules/software/xz/5.2.2/lib -L/beegfs/common/software/2017/modules/software/pcre/8.39/lib -L/beegfs/common/software/2017/modules/software/curl/7.51.0/lib -L/beegfs/common/software/2017/modules/software/rlang/3.3.2/lib64/R/lib"
```

15,1

All

module file

```
#%Module3.2.10#####
proc ModulesHelp { } {
    global version
    puts stderr "\n\tVersion 2.7.12 of python\n"
}
```

Help message shown for this modules with `module help python/2.7.12`

```
module-whatis "Version 2.7.12 of python"
```

Message to be shown by `module whatis python/2.7.12`.

```
# for Tcl script use only
set version "3.2.10"
```

Modules system version.

```
conflict python
```

Conflicts. ie. Do not load if a different python version is already loaded.

```
prepend-path PATH /beegfs/common/software/2017/modules/software/python/2.7.12/bin
prepend-path LD_LIBRARY_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/lib
prepend-path CPATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path C_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path CPLUS_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path_OBJC_INCLUDE_PATH /beegfs/common/software/2017/modules/software/python/2.7.12/include
prepend-path MANPATH /beegfs/common/software/2017/modules/software/python/2.7.12/share/man
prepend-path INFODIR /beegfs/common/software/2017/modules/software/python/2.7.12/share/man
```

Paths prepended when the module is loaded and removed when the module is unloaded.

module file

```
set home $::env(HOME)
```

`\$::env(HOME)` captures the variable HOME from the environment of the user loading the module.

`set home \$::env(HOME)` sets a variable `home` to be used within the module with the result of `\$::env(HOME)`.

```
set pythonuser $home/.python/2.7.12/bin
```

Sets the python user variable to be used within the module.

```
exec /bin/mkdir -p $pythonuser
```

Executes the command `/bin/mkdir -p \$pythonuser` every time the module is loaded.

```
prepend-path PATH $home/.python/2.7.12/bin
```

Prepends `\$home/.python/2.7.12/bin` to the user environment PATH variable.

module file

```
setenv PYTHONUSERBASE $home/.python/2.7.12/
exec /bin/mkdir -p $home/.python/2.7.12/pythonpath/site-packages
```

Please notice `PYTHONUSERBASE`

```
module load gcc/6.3.0 bzip2/1.0.6 xz/5.2.2 ncurses/6.0 libevent/2.0.22 pcre/8.39 curl/7.51.0 freetype/2.8.1
```

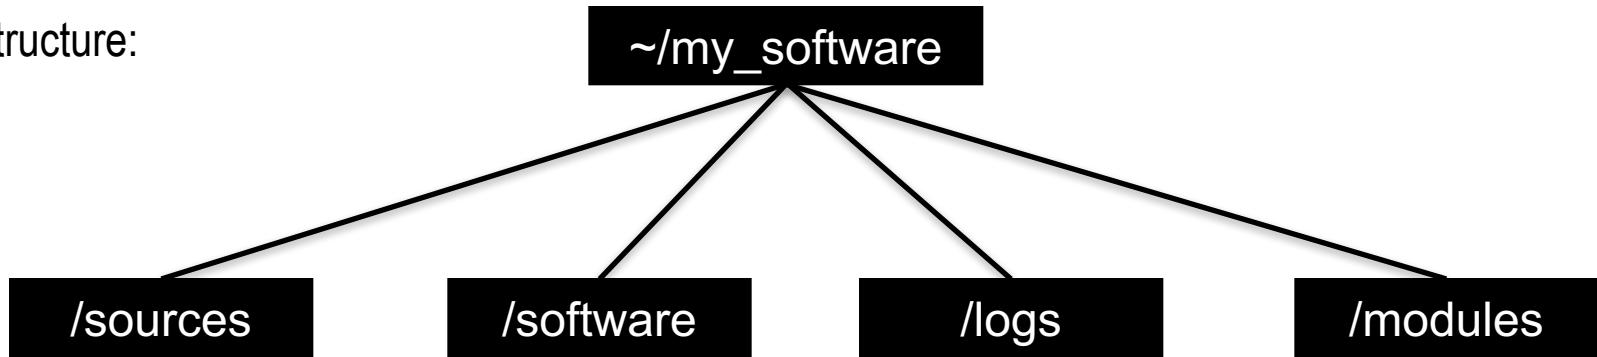
Loads other modules as dependencies when loading the python module.

Installing software without `su` access

Several examples can be obtained from:

https://github.com/mpg-age-bioinformatics/draco_pipelines/blob/master/software/software.sh

Structure:



```
mkdir -p ~/my_software/sources ~/my_software/software \  
~/my_software/logs ~/my_software/modules
```

Installing software without `su` access

Download and edit the script for automatic generation of module files:

```
cd ~/  
wget https://raw.githubusercontent.com/mpg-age-bioinformatics/  
draco_pipelines/master/software/newmod.sh  
  
sed -i 's/\u002fjboucas\u002fmodules/\u002fmy_software/g' newmod.sh  
chmod +x newmod.sh
```

Add the `~/my_software/modules` to the MODULEPATH:

```
export MODULEPATH=$MODULEPATH:~/my_software/modules
```

You might want to add this line to your `~/.bashrc`

Installing software without `su` access

```
cd ~/my_software/sources
wget http://zlib.net/pigz/pigz-2.3.4.tar.gz
tar -zxvf pigz-2.3.4.tar.gz
cd pigz-2.3.4
make
mkdir -p ~/my_software/software/pigz/2.3.4/bin
cp pigz unpigz ~/my_software/software/pigz/2.3.4/bin
~/newmod.sh -s pigz -p ~/my_software/modules -v 2.3.4 -d 2.3.4

# TEST #
module avail
module show pigz/2.3.4
module load pigz/2.3.4
which pigz
```

Installing software without `su` access

```
cd ~/my_software/sources
wget http://zlib.net/zlib-1.2.11.tar.gz
tar -zxvf zlib-1.2.11.tar.gz
cd zlib-1.2.11
make
mkdir -p ~/my_software/software/libz/1.2.11
./configure --prefix=~/my_software/software/libz/1.2.11
make
make install
~/newmod.sh -s libz -p ~/my_software/modules -v 1.2.11 -d 1.2.11
```

Installing software without `su` access

```
echo "#!/bin/bash
cd ~/my_software/sources
wget http://zlib.net/zlib-1.2.11.tar.gz
tar -zxvf zlib-1.2.11.tar.gz
cd zlib-1.2.11
make
mkdir -p ~/my_software/software/libz/1.2.11
./configure --prefix=~/my_software/software/libz/1.2.11
make
make install
~/newmod.sh -s libz -p ~/my_software/modules -v 1.2.11 -d 1.2.11
" > ~/my_software/logs/libz-1.2.11.sh
chmod +x ~/my_software/logs/libz-1.2.11.sh
~/my_software/logs/libz-1.2.11.sh 2>&1 | tee \
~/my_software/logs/libz-1.2.11.log
```

Installing software without `su` access

```
if [ ! -f $MODF/general/tmux/2.3 ]; then
    echo 'tmux-2.3'
    echo '#!/bin/bash'
    module list
    cd $SOUR && \
    wget -O t.tar.gz https://github.com/tmux/tmux/releases/download/2.3/tmux-2.3.ta
r.gz && \
    mv t.tar.gz tmux-2.3.tar.gz && \
    tar -zvxf tmux-2.3.tar.gz && \
    cd tmux-2.3 && \
    mkdir -p $SOFT/tmux/2.3/ && \
    ./configure --prefix=$SOFT/tmux/2.3/ CFLAGS="-I$SOFT/libevent/2.0.22/include -I
$SOFT/ncurses/6.0/include/ncurses" LDFLAGS="-L$SOFT/libevent/2.0.22/lib -L$SOFT/ncu
rses/6.0/lib" && \
    make && make install
    newmod.sh \
    -s tmux \
    -p $MODF/general/ \
    -v 2.3 \
    -d 2.3
    echo "set home $::env(HOME)" >> $MODF/general/tmux/2.3
    echo "exec /bin/mkdir -p \$home/.tmux.socket" >> $MODF/general/tmux/2.3
    echo "module load ncurses/6.0" >> $MODF/general/tmux/2.3
    echo "module load libevent/2.0.22" >> $MODF/general/tmux/2.3
    echo "setenv TMUX_TMPDIR \$home/.tmux.socket" >> $MODF/general/tmux/2.3
    ' > $LOGS/tmux-2.3.sh
    chmod 755 $LOGS/tmux-2.3.sh
    srun -o $LOGS/tmux-2.3.out $LOGS/tmux-2.3.sh
fi
```

Installing R packages

R standard packages:

```
base  datasets  graphics  grDevices  methods  parallel  stats  stats4  
tcltk  tools  utils
```

This are the only packages that should be inside `lib64/R/library`. The `library` folder should have permissions `chmod -R 555 library`. This will enforce the use of the `R_LIBS_USER` variable for installation of new packages.

Alternative (eg. install biomaRt):

```
source("https://bioconductor.org/biocLite.R")  
biocLite("biomaRt", lib=Sys.getenv("R_LIBS_USER"))
```

The `lib` argument can also be used in (eg. install packrat):

```
install.packages("packrat", lib=Sys.getenv("R_LIBS_USER"))
```

Installing python packages

Installing the AGEpy package:

```
module load python  
pip install AGEpy --user
```

Alternative 1:

```
module load python  
git clone https://github.com/mpg-age-bioinformatics/AGEpy  
cd AGEpy  
pip install ../AGEpy --user
```

Alternative 1b (no package control):

```
cd AGEpy  
python setup.py install --user
```

The `--user` argument will make use of the `PYTHONUSERBASE` variable.

Installing perl packages

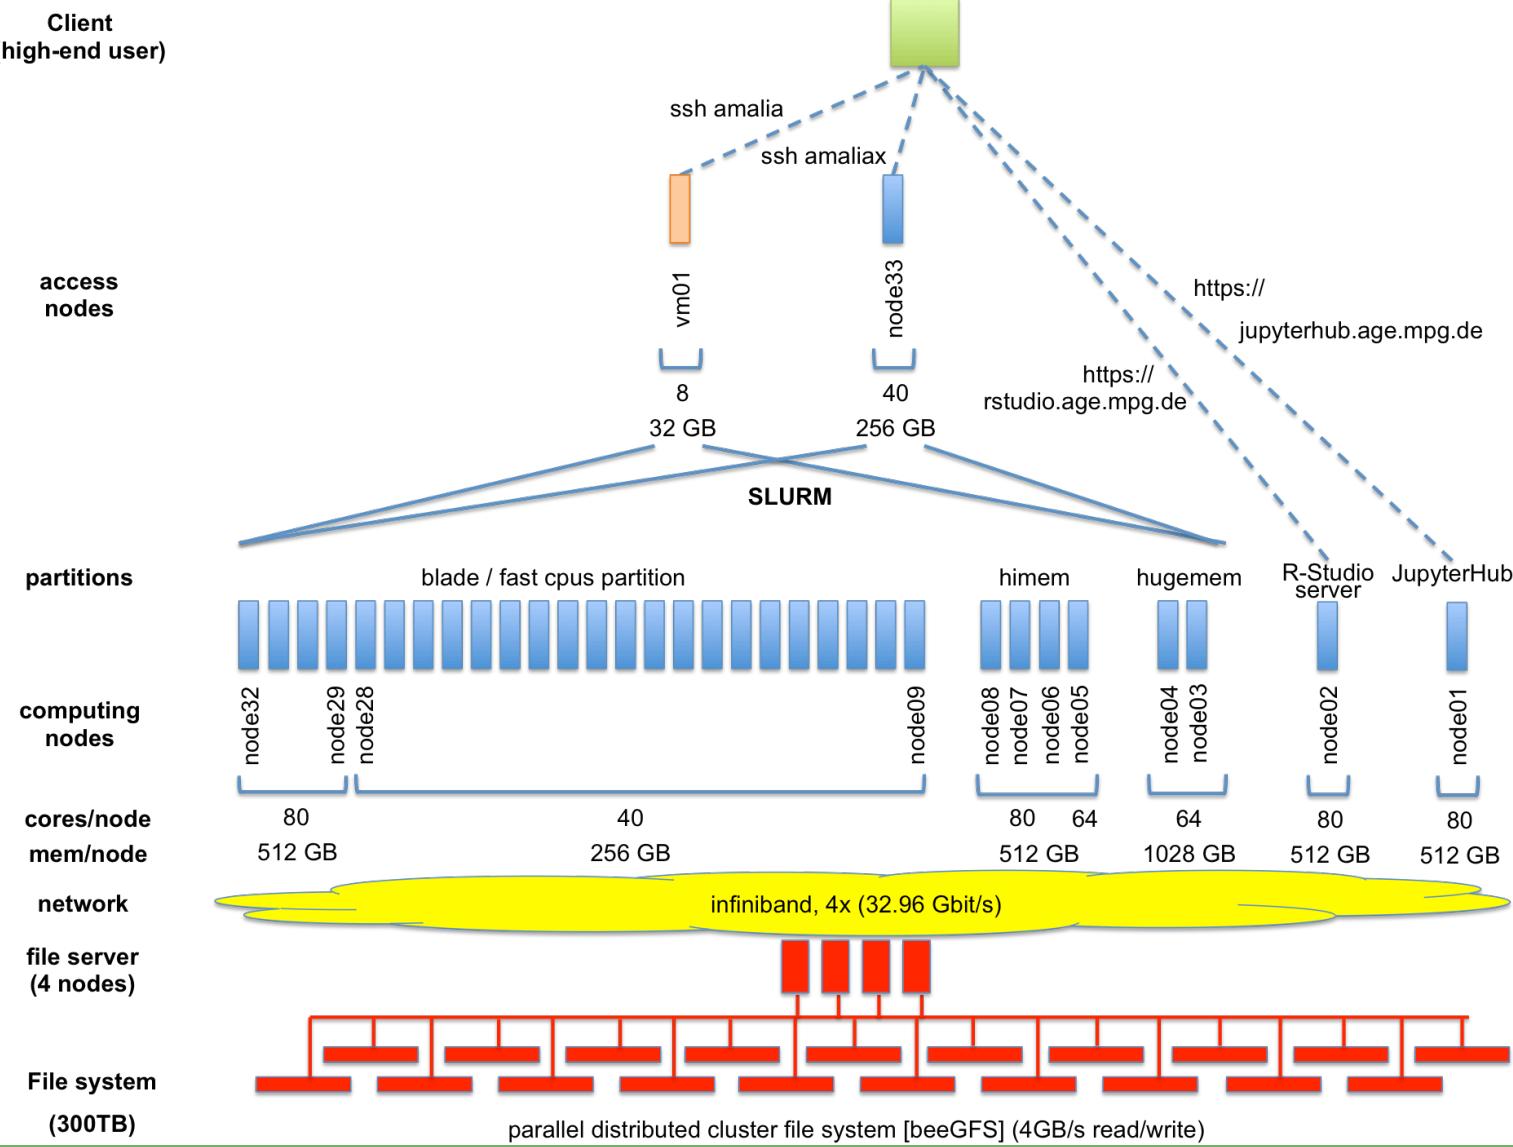
Perl libraries are loaded from the value of `PERL5LIB`. We have added a variable to module perl to ease local installations - `PERLUSER` - check out the content of these two variables.

Example instalation of Log::Log4perl:

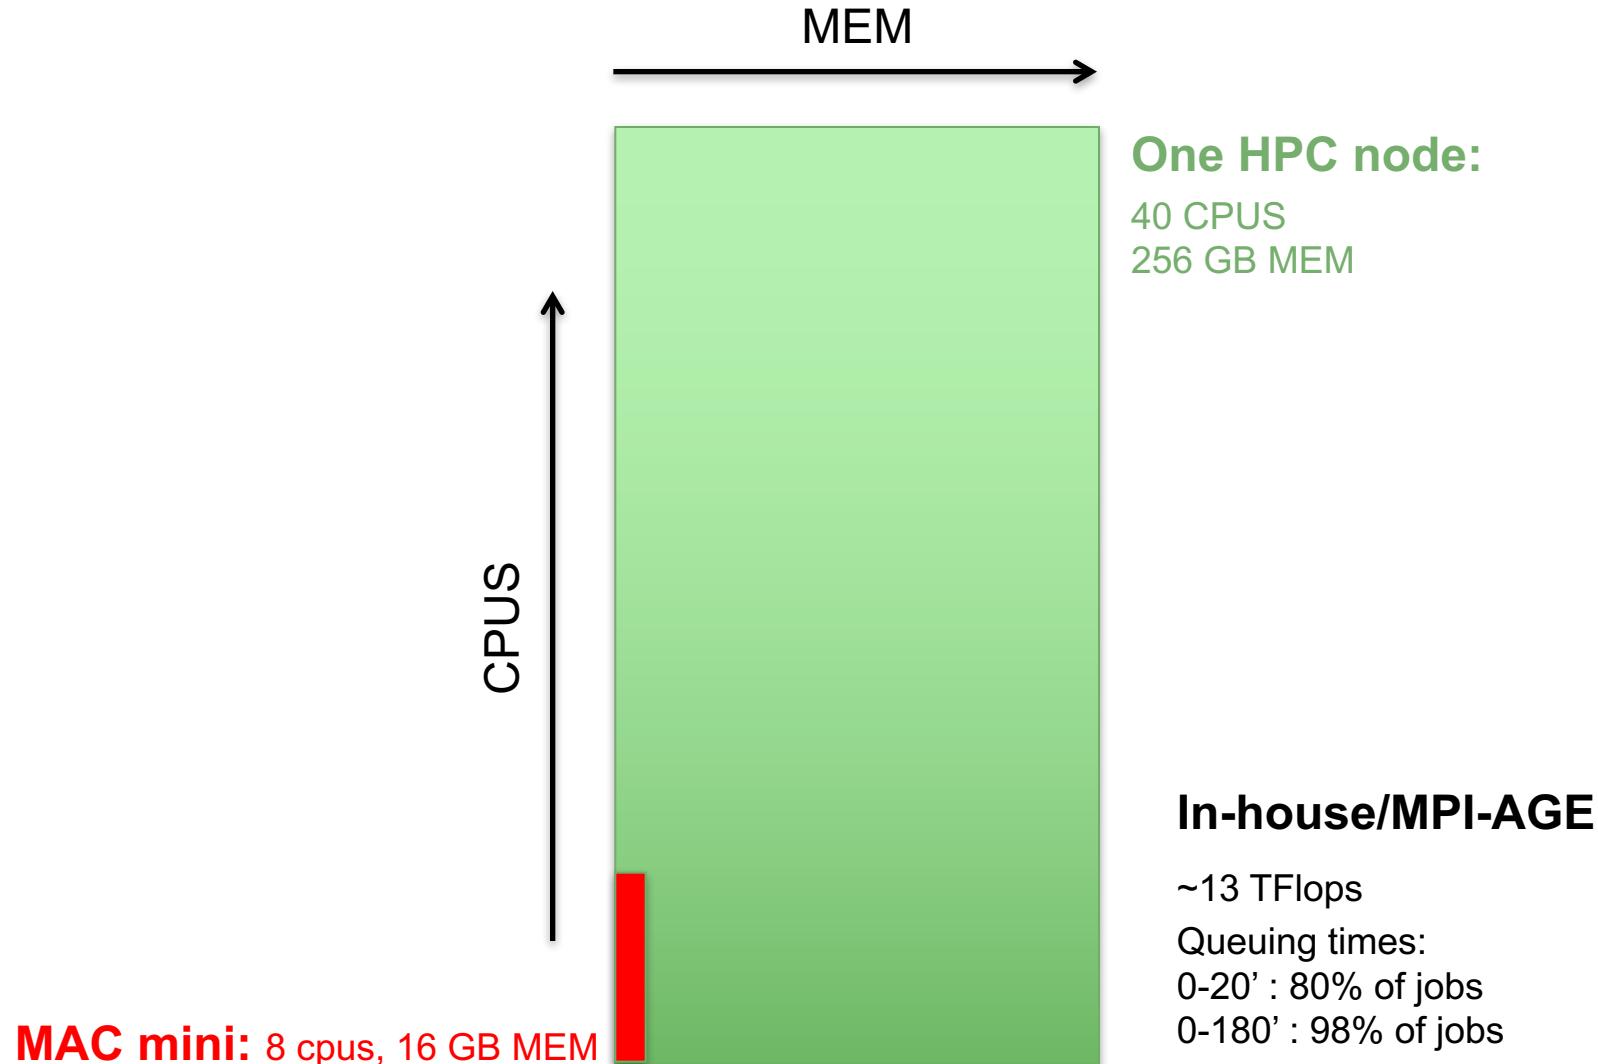
```
module load perl/5.24.1
cpanm Log::Log4perl -l $PERLUSER

# check where it got installed
perldoc -l Log::Log4perl
```

SLURM (Simple Linux Utility for Resource Management)



SLURM: Why an HPC cluster?



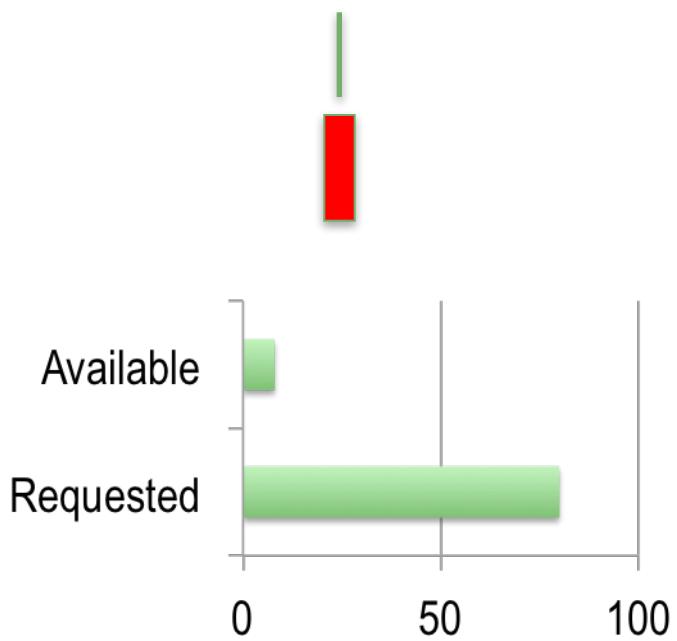
SLURM: Why an HPC cluster?

MAC mini

1 user

4 raw data files

20 cpus / raw file

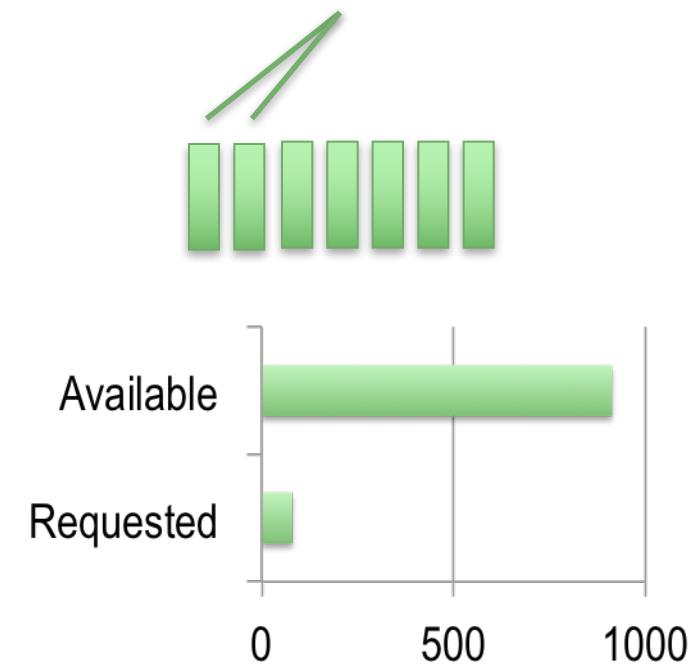


HPC CLUSTER

1 user

4 raw data files

20 cpus / raw file

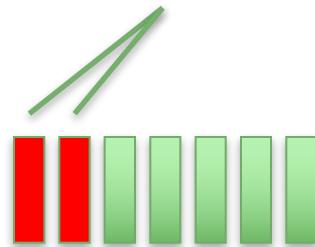


SLURM: Why SLURM?

NO slurm

8 users

20 cpus/user



in use

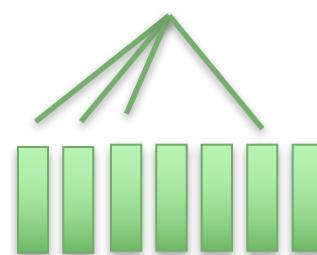
Requested

0 50 100 150 200

with slurm

8 users

20 cpus/user



in use

Requested

0 50 100 150 200

SLURM: Why SLURM?

NO slurm

```
bwa mem sample1.fastq
```

with slurm

```
srun bwa mem sample1.fastq
```

SLURM: How to?

NO slurm

```
bwa mem -T 18 sample1.fastq
```

with slurm

```
srun --cpus-per-task=18 \  
bwa mem -T 18 sample1.fastq
```

SLURM: How to?

NO slurm

```
bwa mem -T 18 sample1.fastq
```

with slurm

```
srun --cpus-per-task=18 \
--mem=64gb \
bwa mem -T 18 sample1.fastq
```

SLURM: How to?

NO slurm

```
bwa mem -T 18 sample1.fastq
```

with slurm

```
srun --cpus-per-task=18 \
--mem=15gb \
--time=5-24 \
bwa mem -T 18 sample1.fastq
```

(5 days and 24 hours = 6 days

alternative: 144:00:00;

maximum = 14 days)

SLURM: How to?

NO slurm

```
bwa mem -T 18 sample1.fastq
```

with slurm

```
srun --cpus-per-task=18 \
--mem=15gb \
--time=5-24 \
-p blade \
bwa mem -T 18 sample1.fastq
```

SLURM: How to?

NO slurm

```
bwa mem -T 18 sample1.fastq
```

with slurm

```
srun --cpus-per-task=18 \
--mem=15gb \
--time=5-24 \
-p blade \
-o slurm_logs/bwa_1.out \
bwa mem -T 18 sample1.fastq
```

SLURM: How to?

NO slurm

```
./align_1.sh
```

```
#!/bin/bash
bwa mem -T 18 sample1.fastq
exit
```

with slurm

```
sbatch --cpus-per-task=18 \
--mem=15gb \
--time=5-24 \
-p blade \
-o slurm_logs/bwa_1.out \
align_1.sh
```

SLURM: scripting

NO slurm

```
#!/bin/bash  
cd ~/project/raw_data  
bwa mem -T 18 sample1.fastq  
exit  
  
../align_1.sh
```

with slurm

```
#!/bin/bash  
#SBATCH --cpus-per-task=18  
#SBATCH --mem=15gb  
#SBATCH --time=5-24  
#SBATCH -p blade  
#SBATCH -o slurm_logs/bwa_1.out  
cd ~/project/raw_data  
bwa mem -T 18 sample1.fastq  
exit  
  
sbatch align_1.sh
```

SLURM with modules system

```
#!/bin/bash

#SBATCH --cpus-per-task=18
#SBATCH --mem=15gb
#SBATCH --time=5-24
#SBATCH -p blade
#SBATCH -o slurm_logs/bwa_1.out

module load bwa
cd ~/project/raw_data
bwa mem -T 18 sample1.fastq
exit
```

SLURM with shifter & the modules system

```
#!/bin/bash

#SBATCH --cpus-per-task=18
#SBATCH --mem=15gb
#SBATCH --time=5-24
#SBATCH -p blade
#SBATCH -o slurm_logs/bwa_1.out

shifter --image=mpgagEBioinformatics/bioinformatics_software:v1.0.1 << SHI
#!/bin/bash
source ~/.bashrc
module load bwa
cd ~/project/raw_data
bwa mem -T 18 sample1.fastq
exit
SHI
```

SLURM: iterating jobs over files

```
#!/bin/bash
cd ~/project/raw_data                                # go to folder containing files

for f in $(ls *.fastq); do echo "#!/bin/bash
cd ~/project/raw_data
bwa mem -T 18 ${f}
rm ~/project/tmp/${f}.sh
" > ~/project/tmp/${f}.sh

chmod 755 ~/project/tmp/${f}.sh
rm ~/project/slurm_logs/${f}.*.out                  # removes pre-existing logs

sbatch --cpus-per-task=18 --mem=15gb \
--time=5-24 -p blade \
-o ~/project/slurm_logs/${f}.%j.out \
~/project/tmp/${f}.sh                                # keeps log with job number

done; exit
```

SLURM: iterating jobs over files

```
#!/bin/bash

cd ~/project/raw_data

for f in $(ls *.fastq);
    do rm ~/project/slurm_logs/${f}.*.out

        sbatch --cpus-per-task=18 --mem=15gb --time=5-24 \
        -p blade -o ~/project/slurm_logs/${f}.%j.out << EOF
#!/bin/bash
cd ~/project/raw_data
bwa mem -T 18 ${f}
EOF

done
exit
```

SLURM: iterating without crashing

```
#!/bin/bash

cd ~/project/raw_data

for f in $(ls *.fastq);
do rm ~/project/slurm_logs/${f}.*.out

while [ `squeue -u Jboucas | wc -l` -gt "500" ];
do echo "sleeping"; sleep 300
done

sbatch --cpus-per-task=18 --mem=15gb --time=5-24 \
-p blade -o ~/project/slurm_logs/${f}.%j.out ~/project/tmp/${f}.sh <<EOF
#!/bin/bash
cd ~/project/raw_data
bwa mem -T 18 ${f}
EOF

done
exit
```

SLURM with shifter

```
#!/bin/bash

cd ~/project/raw_data

for f in $(ls *.fastq);
do   rm ~/project/slurm_logs/${f}.*.out

    sbatch --cpus-per-task=18 --mem=15gb --time=5-24 \
-p blade -o ~/project/slurm_logs/${f}.%j.out ~/project/tmp/${f}.sh <<EOF
#!/bin/bash

shifter -image=mpgabioinformatics/bioinformatics_software:v1.0.1 << SHI
#!/bin/bash

cd ~/project/raw_data

bwa mem -T 18 ${f}

SHI

EOF

done

exit
```

SLURM: other options

```
--mail-type=BEGIN, END,FAIL,REQUEUE,ALL  
  
# Specifies when email is sent to the job owner. The option argument may consist of a  
combination of the allowed mail types  
  
--mail-user=username@age.mpg.de
```

SLURM: sview

The screenshot shows the sview application interface. At the top, there's a menu bar with 'Grab' (highlighted), File, Edit, Capture, Window, Help. To the right of the menu is a status bar showing battery level (100%), time (Thu 17:26), and user (JBoucas). Below the menu is a toolbar with icons for Actions, Options, Query, and Help. A tab bar below the toolbar has tabs for Jobs (selected), Partitions, Reservations, Nodes, and Visible Tabs. The main area is a table displaying job information:

JobID	Partition	UserID	Name	State	Time Running	Node Count	NodeList
165733	blade	wBradshaw	interactive_shell_wBradshaw	RUNNING	7-20:41:54	1	bioinf-node05
166640	blade	wBradshaw	interactive_shell_wBradshaw	RUNNING	6-06:17:45	1	bioinf-node15
166970	blade	MPiechotta	interactive_shell_MPiechotta	RUNNING	5-00:24:27	1	bioinf-node15
167097	blade	MPiechotta	interactive_shell_MPiechotta	RUNNING	4-03:34:43	1	bioinf-node15
167648	blade	RSehlke	rsem-calculate-expression	RUNNING	3-01:49:49	1	bioinf-node09
168275	blade	MPiechotta	interactive_shell_MPiechotta	RUNNING	1-05:59:20	1	bioinf-node09
168353	blade	STempler	interactive_shell_STempler	RUNNING	09:05:15	1	bioinf-node10
168504	himem	JBoucas	go.sh	RUNNING	00:01:14	1	bioinf-node03

A context menu is open over the last row (JobID 168504, Partition himem, UserID JBoucas, Name go.sh). The menu options are: Full Info, Signal, Requeue, Cancel, Suspend/Resume, Edit Job, Partition, Nodes, and Reservation.

SLURM: other useful commands

```
# show the partitions  
  
sinfo  
  
  
# show information on nodes  
  
sinfo -N -O partitionname,nodehost,cpus,cpusload,freemem,memory  
  
  
# requires X forwarding and allows viewing and manipulation of submitted jobs  
  
sview  
  
  
  
# lists running jobs  
  
squeue  
  
  
  
# show the queue for a user  
  
squeue -u <user name>
```

SLURM: other useful commands

```
# cancels job 65673
scancel 65673

# cancels all jobs of user JBoucas
scancel -u JBoucas

# shows detailed resource information on job 43433
scontrol show job 43433

# show information on a partition
scontrol show partition <partition name>

# show information on a node
scontrol show node <node name>
```

SLURM: other useful commands

```
# starts an interactive terminal window on node15
srun -w node15 --pty bash

# submit a job3 after job1 and job2 are successfully ready
job1=$(sbatch --parsable <script1>)
job2=$(sbatch --parsable <script2>)
sbatch -d afterok:${job1}:${job2} <script3>

# attach to a running job and run a command
srun --jobid <JOBID> --pty <command>

# change the partitions of a pending job
scontrol update job <job id> partition=<partition1>,<partition2>,<partition3>
```

Notes:

For current information on RStudio server and JupyterHub please check our cluster first steps page:

https://github.com/mpg-age-bioinformatics/cluster_first_steps

SLURM manual:

http://slurm.schedmd.com/man_index.html

Coming soon:

docker on HPC

END

bioinformatics@age.mpg.de

<https://mpg-age-bioinformatics.github.io>